

Week 10: Data Wrangling in Python

POP77001 Computer Programming for Social Scientists

Tom Paskhalis

14 November 2022

Module website: tinyurl.com/POP77001

Overview

- Numerical analysis in Python
- Tabular data
- Pandas object types
- Working with dataframes in pandas
- Data input and output

Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

```
In [2]: # Subsetting 2nd row, 3rd element  
mat[1][2]
```

```
Out[2]: 6
```

Numerical analysis in Python

- As opposed to other programming languages (Julia, R, MatLab), Python provides very bare bones functionality for numeric analysis.
- E.g. no built-in matrix/array object type, limited mathematical and statistical functions

```
In [1]: # Representing 3x3 matrix with list  
mat = [[1, 2, 3],  
        [4, 5, 6],  
        [7, 8, 9]]
```

```
In [2]: # Subsetting 2nd row, 3rd element  
mat[1][2]
```

```
Out[2]: 6
```

```
In [3]: # Naturally, this representation  
# breaks down rather quickly  
mat * 2
```

```
Out[3]: [[1, 2, 3], [4, 5, 6], [7, 8, 9], [1, 2, 3], [4, 5, 6], [7, 8,  
9]]
```

NumPy - numerical analysis in Python

- NumPy (**N**umeric **P**ython) package provides the basis of numerical computing in Python:
 - multidimensional array
 - mathematical functions for arrays
 - array data I/O
 - linear algebra, RNG, FFT, ...

NumPy - numerical analysis in Python

- NumPy (**N**umeric **P**ython) package provides the basis of numerical computing in Python:
 - multidimensional array
 - mathematical functions for arrays
 - array data I/O
 - linear algebra, RNG, FFT, ...

```
In [4]: # Using 'as' allows to avoid typing full name  
# each time the module is referred to  
import numpy as np
```


NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

```
In [6]: arr[1][2]
```

```
Out[6]: 6
```

NumPy array

- Multidimensional (N) array object (aka ndarray) is a principal container for datasets in Python.
- It is the backbone of data frames, operating behind the scenes

```
In [5]: arr = np.array([[1, 2, 3],  
                        [4, 5, 6],  
                        [7, 8, 9]])
```

```
In [6]: arr[1][2]
```

```
Out[6]: 6
```

```
In [7]: arr * 2
```

```
Out[7]: array([[ 2,  4,  6],  
              [ 8, 10, 12],  
              [14, 16, 18]])
```

Working with arrays

Working with arrays

```
In [8]: # Object type  
        type(arr)
```

```
Out[8]: numpy.ndarray
```

Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

```
In [10]: # Array size  
arr.shape
```

```
Out[10]: (3, 3)
```


Working with arrays

```
In [8]: # Object type  
type(arr)
```

```
Out[8]: numpy.ndarray
```

```
In [9]: # Array dimensionality  
arr.ndim
```

```
Out[9]: 2
```

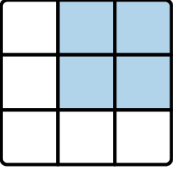
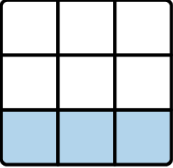
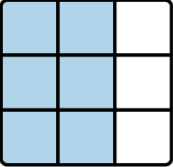
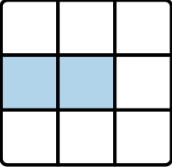
```
In [10]: # Array size  
arr.shape
```

```
Out[10]: (3, 3)
```

```
In [11]: # Calculating summary statistics on array  
# axis indicates the dimension  
# compare to R's `apply(arr, 1, mean)`  
# note that every list within a list  
# is treated as a column (not row)  
arr.mean(axis = 0)
```

```
Out[11]: array([4., 5., 6.])
```

Array indexing and slicing

	Expression	Shape
	<code>arr[:2,1:]</code>	<code>(2,2)</code>
	<code>arr[2]</code> <code>arr[2, :]</code> <code>arr[2:, :]</code>	<code>(3,)</code> <code>(3,)</code> <code>(1,3)</code>
	<code>arr[:, :2]</code>	<code>(3,2)</code>
	<code>arr[1, :2]</code> <code>arr[1:2, :2]</code>	<code>(2,)</code> <code>(1,2)</code>

Source: [Python for Data Analysis](#)

Tidy data

- Tidy data is a specific subset of rectangular data, where:
 - Each variable is in a column
 - Each observation is in a row
 - Each value is in a cell

country	year	cases	population
Afghanistan	1999	18145	19987071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	216766	128062583

variables

country	year	cases	population
Afghanistan	1999	18145	19987071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	216766	128062583

observations

country	year	cases	population
Afghanistan	1999	18145	19987071
Afghanistan	2000	18666	20095360
Brazil	1999	30737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	216766	128062583

values

Source: [R for Data Science](#)

Pandas

- Standard Python library does not have data type for rectangular data
- However, `pandas` library has become the de facto standard for data manipulation
- `pandas` is built upon (and often used in conjunction with) other computational libraries
- E.g. `numpy` (array data type), `scipy` (linear algebra) and `scikit-learn` (machine learning)

Pandas

- Standard Python library does not have data type for rectangular data
- However, `pandas` library has become the de facto standard for data manipulation
- `pandas` is built upon (and often used in conjunction with) other computational libraries
- E.g. `numpy` (array data type), `scipy` (linear algebra) and `scikit-learn` (machine learning)

```
In [12]: import pandas as pd
```

Core `pandas` object types

- *Series* - one-dimensional sequence of values
- *DataFrame* - (typically) two-dimensional rectangular table

Series

- *Series* is a one-dimensional array-like object

Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0      150.0  
         1      120.0  
         2     3000.0  
         dtype: float64
```


Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0      150.0  
         1      120.0  
         2     3000.0  
         dtype: float64
```

```
In [14]: sr1[0] # Slicing is similar to standard Python objects
```

```
Out[14]: 150.0
```

Series

- *Series* is a one-dimensional array-like object

```
In [13]: sr1 = pd.Series([150.0, 120.0, 3000.0])  
sr1
```

```
Out[13]: 0      150.0  
         1      120.0  
         2     3000.0  
         dtype: float64
```

```
In [14]: sr1[0] # Slicing is similar to standard Python objects
```

```
Out[14]: 150.0
```

```
In [15]: sr1[sr1 > 200] # But subsetting is also available
```

```
Out[15]: 2      3000.0  
         dtype: float64
```

Indexing in Series

- Another way to think about Series is as a ordered dictionary

Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)  
sr2
```

```
Out[17]: apple      150.0  
banana      120.0  
watermelon  3000.0  
dtype: float64
```

Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)
sr2
```

```
Out[17]: apple      150.0
         banana     120.0
         watermelon 3000.0
         dtype: float64
```

```
In [18]: sr2[0] # Recall that this slicing would be impossible for standard dict
```

```
Out[18]: 150.0
```

Indexing in Series

- Another way to think about Series is as a ordered dictionary

```
In [16]: d = {'apple': 150.0, 'banana': 120.0, 'watermelon': 3000.0}
```

```
In [17]: sr2 = pd.Series(d)
sr2
```

```
Out[17]: apple          150.0
         banana         120.0
         watermelon    3000.0
         dtype: float64
```

```
In [18]: sr2[0] # Recall that this slicing would be impossible for standard dict
```

```
Out[18]: 150.0
```

```
In [19]: sr2.index # Sequence of labels is converted into an Index object
```

```
Out[19]: Index(['apple', 'banana', 'watermelon'], dtype='object')
```

DataFrame - the workhorse of data analysis

- *DataFrame* is a rectangular table of data

DataFrame - the workhorse of data analysis

- *DataFrame* is a rectangular table of data

```
In [20]: data = {'fruit': ['apple', 'banana', 'watermelon'], # DataFrame can be
                'weight': [150.0, 120.0, 3000.0],           # a dict of equal-l
                'berry': [False, True, True]}
df = pd.DataFrame(data)
df
```

```
Out[20]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True
2	watermelon	3000.0	True

Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

```
In [21]: df.iloc[0] # First row
```

```
Out[21]: fruit      apple  
weight    150.0  
berry      False  
Name: 0, dtype: object
```

Indexing in DataFrame

- DataFrame has both row and column indices
- `DataFrame.loc()` provides method for *label* location
- `DataFrame.iloc()` provides method for *index* location

```
In [21]: df.iloc[0] # First row
```

```
Out[21]: fruit      apple
         weight    150.0
         berry     False
         Name: 0, dtype: object
```

```
In [22]: df.iloc[:,0] # First column
```

```
Out[22]: 0      apple
         1     banana
         2  watermelon
         Name: fruit, dtype: object
```

Summary of indexing in DataFrame

Expression	Selection Operation
<code>df[val]</code>	Column or sequence of columns +convenience (e.g. slice)
<code>df.loc[lab_i]</code>	Row or subset of rows by label
<code>df.loc[:, lab_j]</code>	Column or subset of columns by label
<code>df.loc[lab_i, lab_j]</code>	Both rows and columns by label
<code>df.iloc[i]</code>	Row or subset of rows by integer position
<code>df.iloc[:, j]</code>	Column or subset of columns by integer position
<code>df.iloc[i, j]</code>	Both rows and columns by integer position
<code>df.at[lab_i, lab_j]</code>	Single scalar value by row and column label
<code>df.iat[i, j]</code>	Single scalar value by row and column integer position

Extra: [Pandas documentation on indexing](#)

Subsetting in DataFrame

Subsetting in DataFrame

In [23]: `df.iloc[:2]` *# Select the first two rows (with convenience shortcut for*

Out[23]:

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [25]: df.loc[:, ['fruit', 'berry']] # Select the columns 'fruit' and 'berry'
```

```
Out[25]:
```

	fruit	berry
0	apple	False
1	banana	True
2	watermelon	True

Subsetting in DataFrame

```
In [23]: df.iloc[:2] # Select the first two rows (with convenience shortcut for
```

```
Out[23]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [24]: df[:2] # Shortcut
```

```
Out[24]:
```

	fruit	weight	berry
0	apple	150.0	False
1	banana	120.0	True

```
In [25]: df.loc[:, ['fruit', 'berry']] # Select the columns 'fruit' and 'berry'
```

```
Out[25]:
```

	fruit	berry
0	apple	False
1	banana	True
2	watermelon	True

```
In [26]: df[['fruit', 'berry']] # Shortcut
```

Out[26]:

	fruit	berry
0	apple	False
1	banana	True
2	watermelon	True

Columns in DataFrame

Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns (index object)
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns (index object)
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subsettable
```

```
Out[28]: 'fruit'
```

Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns (index object)
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subsettable
```

```
Out[28]: 'fruit'
```

```
In [29]: df.columns.str.startswith('fr') # As column names are strings, we can a
```

```
Out[29]: array([ True, False, False])
```

Columns in DataFrame

```
In [27]: df.columns # Retrieve the names of all columns (index object)
```

```
Out[27]: Index(['fruit', 'weight', 'berry'], dtype='object')
```

```
In [28]: df.columns[0] # This Index object is subsettable
```

```
Out[28]: 'fruit'
```

```
In [29]: df.columns.str.startswith('fr') # As column names are strings, we can a
```

```
Out[29]: array([ True, False, False])
```

```
In [30]: df.iloc[:,df.columns.str.startswith('fr')] # This is helpful with more
```

```
Out[30]:
```

	fruit
0	apple
1	banana
2	watermelon

Filtering in DataFrame

Filtering in DataFrame

```
In [31]: # Select rows where fruits are not berries  
df[df.loc[:, 'berry'] == False]
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

Filtering in DataFrame

```
In [31]: # Select rows where fruits are not berries  
df[df.loc[:, 'berry'] == False]
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [32]: # The same can be achieved with more concise syntax  
df[df['berry'] == False]
```

```
Out[32]:
```

	fruit	weight	berry
0	apple	150.0	False

Filtering in DataFrame

```
In [31]: # Select rows where fruits are not berries  
df[df.loc[:, 'berry'] == False]
```

```
Out[31]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [32]: # The same can be achieved with more concise syntax  
df[df['berry'] == False]
```

```
Out[32]:
```

	fruit	weight	berry
0	apple	150.0	False

```
In [33]: # Create new dataset with rows where weight is higher than 200  
weight200 = df[df['weight'] > 200]  
weight200
```

```
Out[33]:
```

	fruit	weight	berry
2	watermelon	3000.0	True

Manipulating columns in DataFrame

Manipulating columns in DataFrame

```
In [34]: df = df.rename(columns = {'weight': 'weight_g'}) # Columns can be renamed
```

Manipulating columns in DataFrame

```
In [34]: df = df.rename(columns = {'weight': 'weight_g'}) # Columns can be renamed
```

```
In [35]: df
```

```
Out[35]:
```

	fruit	weight_g	berry
0	apple	150.0	False
1	banana	120.0	True
2	watermelon	3000.0	True

Manipulating columns in DataFrame

```
In [34]: df = df.rename(columns = {'weight': 'weight_g'}) # Columns can be renamed
```

```
In [35]: df
```

```
Out[35]:
```

	fruit	weight_g	berry
0	apple	150.0	False
1	banana	120.0	True
2	watermelon	3000.0	True

```
In [36]: df['weight_oz'] = 0 # Columns can be added or modified by assignment
```

Manipulating columns in DataFrame

```
In [34]: df = df.rename(columns = {'weight': 'weight_g'}) # Columns can be renamed
```

```
In [35]: df
```

```
Out[35]:
```

	fruit	weight_g	berry
0	apple	150.0	False
1	banana	120.0	True
2	watermelon	3000.0	True

```
In [36]: df['weight_oz'] = 0 # Columns can be added or modified by assignment
```

```
In [37]: df
```

```
Out[37]:
```

	fruit	weight_g	berry	weight_oz
0	apple	150.0	False	0
1	banana	120.0	True	0
2	watermelon	3000.0	True	0

Manipulating columns in DataFrame continued

Manipulating columns in DataFrame continued

```
In [38]: df['weight_oz'] = df['weight_g'] * 0.04
```

Manipulating columns in DataFrame continued

```
In [38]: df['weight_oz'] = df['weight_g'] * 0.04
```

```
In [39]: df
```

```
Out[39]:
```

	fruit	weight_g	berry	weight_oz
0	apple	150.0	False	6.0
1	banana	120.0	True	4.8
2	watermelon	3000.0	True	120.0

Manipulating columns in DataFrame continued

```
In [38]: df['weight_oz'] = df['weight_g'] * 0.04
```

```
In [39]: df
```

```
Out[39]:
```

	fruit	weight_g	berry	weight_oz
0	apple	150.0	False	6.0
1	banana	120.0	True	4.8
2	watermelon	3000.0	True	120.0

```
In [40]: df['weight_oz'] = round(df['weight_oz'], 1)
```

Manipulating columns in DataFrame continued

```
In [38]: df['weight_oz'] = df['weight_g'] * 0.04
```

```
In [39]: df
```

```
Out[39]:
```

	fruit	weight_g	berry	weight_oz
0	apple	150.0	False	6.0
1	banana	120.0	True	4.8
2	watermelon	3000.0	True	120.0

```
In [40]: df['weight_oz'] = round(df['weight_oz'], 1)
```

```
In [41]: df
```

```
Out[41]:
```

	fruit	weight_g	berry	weight_oz
0	apple	150.0	False	6.0
1	banana	120.0	True	4.8
2	watermelon	3000.0	True	120.0

Variable transformation

- Lambda functions can be used to transform data with `map ()` method

Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [42]: df['fruit'].map(lambda x: x.upper())
```

```
Out[42]: 0      APPLE  
1     BANANA  
2  WATERMELON  
Name: fruit, dtype: object
```

Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [42]: df['fruit'].map(lambda x: x.upper())
```

```
Out[42]: 0      APPLE  
         1     BANANA  
         2  WATERMELON  
         Name: fruit, dtype: object
```

```
In [43]: transform = lambda x: x.capitalize()
```

Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [42]: df['fruit'].map(lambda x: x.upper())
```

```
Out[42]: 0      APPLE  
1     BANANA  
2  WATERMELON  
Name: fruit, dtype: object
```

```
In [43]: transform = lambda x: x.capitalize()
```

```
In [44]: transformed = df['fruit'].map(transform)
```

Variable transformation

- Lambda functions can be used to transform data with `map()` method

```
In [42]: df['fruit'].map(lambda x: x.upper())
```

```
Out[42]: 0      APPLE  
1     BANANA  
2  WATERMELON  
Name: fruit, dtype: object
```

```
In [43]: transform = lambda x: x.capitalize()
```

```
In [44]: transformed = df['fruit'].map(transform)
```

```
In [45]: transformed
```

```
Out[45]: 0      Apple  
1     Banana  
2  Watermelon  
Name: fruit, dtype: object
```

File object

- File object in Python provides the main interface to external files
- In contrast to other core types, file objects are created not with a literal,
- But with a function, `open ()` :

```
<variable_name> = open(<filepath>, <mode>)
```


Data input and output

- Modes of file objects allow to:
 - (r)ead a file (default)
 - (w)rite an object to a file
 - e(x)clusively create, failing if a file exists
 - (a)ppend to a file
- You can r+ mode if you need to read and write to file

Data output example

Data output example

```
In [46]: f = open('../temp/test.txt', 'w') # Create a new file object in write mode
```

Data output example

```
In [46]: f = open('../temp/test.txt', 'w') # Create a new file object in write mode
```

```
In [47]: f.write('This is a test file.') # Write a string of characters to it
```

```
Out[47]: 20
```

Data output example

```
In [46]: f = open('../temp/test.txt', 'w') # Create a new file object in write mode
```

```
In [47]: f.write('This is a test file.') # Write a string of characters to it
```

```
Out[47]: 20
```

```
In [48]: f.close() # Flush output buffers to disk and close the connection
```

Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

```
In [49]: with open('../temp/test.txt', 'r') as f: # Note that we use 'r' mode for reading
          text = f.read()
```

Data input example

- To avoid keeping track of open file connections, `with` statement can be used

Extra: [Python documentation on with statement](#)

```
In [49]: with open('../temp/test.txt', 'r') as f: # Note that we use 'r' mode for reading
          text = f.read()
```

```
In [50]: text
```

```
Out[50]: 'This is a test file.'
```


Reading and writing data in pandas

- pandas provides high-level methods that takes care of file connections
- These methods all follow the same read_<format> and to_<format> name patterns
- CSV (comma-separated value) files are the standard of interoperability

```
<variable_name> = pd.read_<format>(<filepath>)
```

```
<variable_name>.to_<format>(<filepath>)
```

Reading data in `pandas` example

- We will use the data from [Kaggle 2021 Machine Learning and Data Science Survey](#)
- For more information you can read the [executive summary](#)
- Or explore the [winning Python Jupyter Notebooks](#)

Reading data in `pandas` example

- We will use the data from [Kaggle 2021 Machine Learning and Data Science Survey](#)
- For more information you can read the [executive summary](#)
- Or explore the [winning Python Jupyter Notebooks](#)

```
In [51]: # We specify that we want to combine first two rows as a header
kaggle2021 = pd.read_csv(
    '../data/kaggle_survey_2021_responses.csv',
    header = [0,1]
)
```

```
/tmp/ipykernel_215835/1877677354.py:2: DtypeWarning: Columns (1
95,201) have mixed types. Specify dtype option on import or set
low_memory=False.
kaggle2021 = pd.read_csv(
```

Visual data inspection

Visual data inspection

```
In [52]: kaggle2021.head() # Returns the top n (n=5 default) rows
```

Out[52]:							
Time from Start to Finish (seconds)		Q1	Q2	Q3	Q4	Q5	
Duration (in seconds)	What is your age (# years)?	What is your gender? - Selected Choice	In which country do you currently reside?	What is the highest level of formal education that you have attained or plan to attain within the next 2 years?	Select the title most similar to your current role (or most recent title if retired): - Selected Choice	For how many years have you written this program	
0	910	50-54	Man	India	Bachelor's degree	Other	5-1

1	784	50-54	Man	Indonesia	Master's degree	Program/Project Manager	20
2	924	22-24	Man	Pakistan	Master's degree	Software Engineer	1-
3	575	45-49	Man	Mexico	Doctoral degree	Research Scientist	20
4	781	45-49	Man	India	Doctoral degree	Other	<

5 rows × 369 columns



Visual data inspection continued

Visual data inspection continued

In [53]: `kaggle2021.tail()` *# Returns the bottom n (n=5 default) rows*

Out[53]:

	Time from Start to Finish (seconds)	Q1	Q2	Q3	Q4	Q5	
Duration (in seconds)	What is your age (# years)?	What is your gender? - Selected Choice	In which country do you currently reside?	What is the highest level of formal education that you have attained or plan to attain within the next 2 years?	Select the title most similar to your current role (or most recent title if retired): - Selected Choice	For how many years you writing a programn	
25968	1756	30-34	Man	Egypt	Bachelor's degree	Data Analyst	1-3

25969	253	22-24	Man	China	Master's degree	Student	1-3
25970	494	50-54	Man	Sweden	Doctoral degree	Research Scientist	I have written
25971	277	45-49	Man	United States of America	Master's degree	Data Scientist	5-10
25972	255	18-21	Man	India	Bachelor's degree	Business Analyst	I have written

5 rows × 369 columns



Reading in other (non-`.csv`) data files

- Pandas can read in file other than `.csv` (comma-separated value)
- Common cases include STATA `.dta`, SPSS `.sav` and SAS `.sas`
- Use `pd.read_stata(path)`, `pd.read_spss(path)` and `pd.read_sas(path)`
- Check [here](#) for more examples

Writing data out in Python

- Note that when writing data out we start with the object name storing the dataset
- I.e. `df.to_csv(path)` as opposed to `df = pd.read_csv(path)`
- Pandas can also write out into other data formats
- E.g. `df.to_excel(path)` , `df.to_stata(path)`

Writing data out in Python

- Note that when writing data out we start with the object name storing the dataset
- I.e. `df.to_csv(path)` as opposed to `df = pd.read_csv(path)`
- Pandas can also write out into other data formats
- E.g. `df.to_excel(path)`, `df.to_stata(path)`

```
In [54]: kaggle2021.to_csv('../temp/kaggle2021.csv')
```

Summarizing numeric variables

- DataFrame methods in pandas can automatically handle (exclude) missing data (NaN)

Summarizing numeric variables

- DataFrame methods in pandas can automatically handle (exclude) missing data (NaN)

```
In [55]: kaggle2021.describe() # DataFrame.describe() provides an range of summary statistics
```

```
Out[55]:
```

Time from Start to Finish (seconds)	Q30_B_Part_1	Q30_B_Part_2	Q30_B_Part_3	Q30_B_Part
<hr/>				

	Duration (in seconds)	In the next 2 years, do you hope to become more familiar with any of these specific data storage products? (Select all that apply) - Microsoft Azure Data Lake Storage	In the next 2 years, do you hope to become more familiar with any of these specific data storage products? (Select all that apply) - Microsoft Azure Disk Storage	In the next 2 years, do you hope to become more familiar with any of these specific data storage products? (Select all that apply) - Amazon Simple Storage Service (S3)	In the next 2 years, do you hope to become more familiar with any of these specific data storage products? (Select all that apply) - Amazon Elastic File System (EFS)
count	2.597300e+04	0.0	0.0	0.0	0.0
mean	1.105466e+04	NaN	NaN	NaN	NaN
std	1.014716e+05	NaN	NaN	NaN	NaN
min	1.200000e+02	NaN	NaN	NaN	NaN
25%	4.430000e+02	NaN	NaN	NaN	NaN
50%	6.560000e+02	NaN	NaN	NaN	NaN
75%	1.038000e+03	NaN	NaN	NaN	NaN
max	2.488653e+06	NaN	NaN	NaN	NaN

Methods for summarizing numeric variables

Methods for summarizing numeric variables

```
In [56]: kaggle2021.iloc[:,0].mean() # Rather than using describe(), we can apply
```

```
Out[56]: 11054.66492126439
```

Methods for summarizing numeric variables

```
In [56]: kaggle2021.iloc[:,0].mean() # Rather than using describe(), we can apply
```

```
Out[56]: 11054.66492126439
```

```
In [57]: kaggle2021.iloc[:,0].median() # Median
```

```
Out[57]: 656.0
```

Methods for summarizing numeric variables

```
In [56]: kaggle2021.iloc[:,0].mean() # Rather than using describe(), we can apply
```

```
Out[56]: 11054.66492126439
```

```
In [57]: kaggle2021.iloc[:,0].median() # Median
```

```
Out[57]: 656.0
```

```
In [58]: kaggle2021.iloc[:,0].std() # Standard deviation
```

```
Out[58]: 101471.6221245172
```

Methods for summarizing numeric variables

```
In [56]: kaggle2021.iloc[:,0].mean() # Rather than using describe(), we can apply
```

```
Out[56]: 11054.66492126439
```

```
In [57]: kaggle2021.iloc[:,0].median() # Median
```

```
Out[57]: 656.0
```

```
In [58]: kaggle2021.iloc[:,0].std() # Standard deviation
```

```
Out[58]: 101471.6221245172
```

```
In [59]: import statistics ## We don't have to rely only on methods provided by  
statistics.stdev(kaggle2021.iloc[:,0])
```

```
Out[59]: 101471.6221245172
```

Summarizing categorical variables

Summarizing categorical variables

In [60]: `kaggle2021.describe(include = 'all')` *# Adding include = 'all' tells pandas to include all variables*

Out[60]:

	Time from Start to Finish (seconds)	Q1	Q2	Q3	Q4	Q5
	Duration (in seconds)	What is your age (# years)?	What is your gender? - Selected Choice	In which country do you currently reside?	What is the highest level of formal education that you have attained or plan to attain within the next 2 years?	Select the title most similar to your current role (or most recent title if retired): - Selected Choice
count	2.597300e+04	25973	25973	25973	25973	25973
unique	NaN	11	5	66	7	15

	top	NaN	25-29	Man	India	Master's degree	Student
	freq	NaN	4931	20598	7434	10132	6804
	mean	1.105466e+04	NaN	NaN	NaN	NaN	NaN
	std	1.014716e+05	NaN	NaN	NaN	NaN	NaN
	min	1.200000e+02	NaN	NaN	NaN	NaN	NaN
	25%	4.430000e+02	NaN	NaN	NaN	NaN	NaN
	50%	6.560000e+02	NaN	NaN	NaN	NaN	NaN
	75%	1.038000e+03	NaN	NaN	NaN	NaN	NaN
	max	2.488653e+06	NaN	NaN	NaN	NaN	NaN

11 rows × 369 columns



Methods for summarizing categorical variables

Methods for summarizing categorical variables

```
In [61]: kaggle2021.iloc[:,2].mode() # Mode, most frequent value
```

```
Out[61]: 0      Man  
Name: (Q2, What is your gender? - Selected Choice), dtype: object
```

Methods for summarizing categorical variables

```
In [61]: kaggle2021.iloc[:,2].mode() # Mode, most frequent value
```

```
Out[61]: 0      Man  
         Name: (Q2, What is your gender? - Selected Choice), dtype: object
```

```
In [62]: kaggle2021.iloc[:,2].value_counts() # Counts of unique values
```

```
Out[62]: Man                20598  
         Woman              4890  
         Prefer not to say    355  
         Nonbinary           88  
         Prefer to self-describe 42  
         Name: (Q2, What is your gender? - Selected Choice), dtype: int64
```

Methods for summarizing categorical variables

```
In [61]: kaggle2021.iloc[:,2].mode() # Mode, most frequent value
```

```
Out[61]: 0    Man
          Name: (Q2, What is your gender? - Selected Choice), dtype: object
```

```
In [62]: kaggle2021.iloc[:,2].value_counts() # Counts of unique values
```

```
Out[62]: Man                20598
          Woman              4890
          Prefer not to say    355
          Nonbinary           88
          Prefer to self-describe 42
          Name: (Q2, What is your gender? - Selected Choice), dtype: int64
```

```
In [63]: kaggle2021.iloc[:,2].value_counts(normalize = True) # We can further normalize
```

```
Out[63]: Man                0.793054
          Woman              0.188272
          Prefer not to say    0.013668
          Nonbinary           0.003388
          Prefer to self-describe 0.001617
```

Summary of descriptive statistics methods

Method	Numeric	Categorical	Description
count	yes	yes	Number of non-NA observations
value_counts	yes	yes	Number of unique observations by value
describe	yes	yes	Set of summary statistics for Series/DataFrame
min, max	yes	yes (caution)	Minimum and maximum values
quantile	yes	no	Sample quantile ranging from 0 to 1
sum	yes	yes (caution)	Sum of values
prod	yes	no	Product of values
mean	yes	no	Mean
median	yes	no	Median (50% quantile)
var	yes	no	Sample variance
std	yes	no	Sample standard deviation
skew	yes	no	Sample skewness (third moment)
kurt	yes	no	Sample kurtosis (fourth moment)

Next

- Tutorial: Reading/writing and reshaping data in Python
- Next week: Classes and Object-oriented programming